

Introduction to Machine Learning Engineering

Model governance for risk management

Musa Baloyi

August 14, 2018

Table of contents

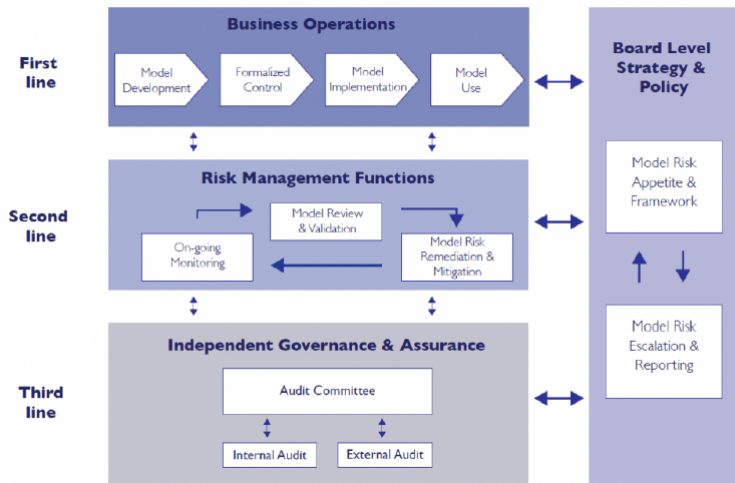
- ▶ Model governance for risk management
 - ▶ Model risk
 - ▶ Model management
 - ▶ Model governance
 - ▶ Application performance management
- ▶ Machine learning engineering
 - ▶ The Big Data Platform
 - ▶ The Elastic Stack
 - ▶ Data generation
 - ▶ Data movement
 - ▶ Architectural guidelines
- ▶ Getting started guide

Model governance for risk management

- ▶ Logging
- ▶ Monitoring
- ▶ Metrics
- ▶ APM
- ▶ Model Governance
- ▶ Model Risk
- ▶ Risk Management
- ▶ Fit for Purpose

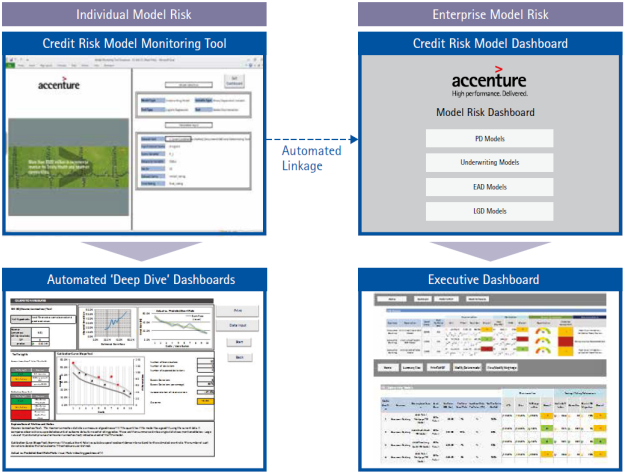
Model governance

8 Key elements for a solid Model Governance framework



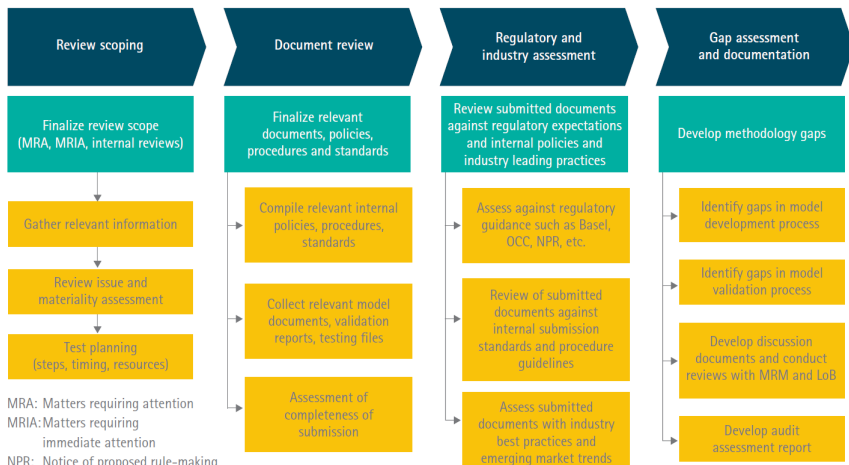
Model governance

Figure 5: Accenture Credit Risk Model Monitoring Suite



Model governance

Figure 4. Audit Review Framework



MRA: Matters requiring attention

MRIA: Matters requiring immediate attention

NPR: Notice of proposed rule-making

Source: Accenture, August 2015

Model governance

Board engagement and communication

Clearly stated
validation
objective,
scope,
framework

Single
internal model
validation
report

Company
specific
confirmation
statements

Use test
questions

Key graphics
and metrics

Standardised
validation test
schedules


Understanding
how external
resources can
provide support

Regular
reporting
and/or agenda
item

The Elastic Stack

- ▶ The Elastic Stack
- ▶ Elasticsearch
- ▶ Kibana
- ▶ ES for Hadoop
- ▶ Architecture
- ▶ Kafka vs Beats
- ▶ Logstash
- ▶ Logging
- ▶ Monitoring
- ▶ Visualisation
- ▶ X-Pack and Machine Learning

The Elastic Stack



Installation guidelines

- ▶ Installing from source
- ▶ Using a package manager
- ▶ sebp/elk docker container
- ▶ elastic/stack-docker docker container
- ▶ Elastic Team SIT
- ▶ Elastic Team SLAM

Elastic Stack demo

- ▶ `sudo docker pull sebp/elk`
- ▶ `sudo sysctl -w vm.max_map_count=300000`
- ▶ `sudo docker run -p 5601:5601 -p 9200:9200 -p 5044:5044 -it --name elk sebp/elk`
- ▶ Elasticsearch is running on `http://localhost:9200`
- ▶ Kibana is running on `http://localhost:5601`
- ▶ Logstash started at 5044

Elasticsearch demo

- ▶ `curl -XGET 'localhost:9200/_cat/health?v&pretty'`
- ▶ `curl -XGET 'localhost:9200/_cat/nodes?v&pretty'`
- ▶ `curl -XGET 'localhost:9200/_cat/indices?v&pretty'`
- ▶ Create index

```
musa@musa-VirtualBox:~$ curl -XPUT 'localhost:9200/rta-all-models2?pretty' -H 'Content-type: application/json' -d '{
  "mappings": {
    "log": {
      "properties": {
        "env_name": {"type": "text"},
        "submit_status": {"type": "text"},
        "model_name": {"type": "text"},
        "tot_runtime": {"type": "float"}
      }
    }
  }
}'
{
  "acknowledged" : true,
  "shards_acknowledged" : true,
  "index" : "rta-all-models2"
}
```

Elasticsearch demo

- ▶ `curl -XGET 'localhost:9200/_cat/indices?v&pretty'`

```
musa@musa-VirtualBox:~$ curl -XGET 'localhost:9200/_cat/indices?v&pretty'
```

| health | status | index | uuid | pri | rep | docs.count | docs.deleted | store.size | pri.store.size |
|--------|--------|---------------------|------------------------|-----|-----|------------|--------------|------------|----------------|
| yellow | open | logstash-2015.05.18 | sKtfBC7HSMWBBgfUcPUidg | 5 | 1 | 0 | 0 | 1.1kb | 1.1kb |
| yellow | open | rta-all-models | MkgSfii6QrytEepm6uBw4w | 5 | 1 | 0 | 0 | 1.1kb | 1.1kb |
| yellow | open | rta-all-models2 | lIp83IE7Qh0FRom6UJsYqg | 5 | 1 | 0 | 0 | 1.1kb | 1.1kb |


- ▶ `curl -H 'Content-Type: application/x-ndjson' -XPOST 'localhost:9200/_bulk?pretty' --data-binary @rta_2018-03-13T08 37 47.143254.json`

Elasticsearch clients



Kibana demo

- Access Kibana to see created indices and data



Kibana demo

► Create index pattern in Kibana

★ rta-all-models*



This page lists every field in the **rta-all-models*** index and the field's associated core type as recorded by Elasticsearch. While this list allows you to view the core type of each field, changing field types must be done using Elasticsearch's [Mapping API](#).

fields (9)

scripted fields (0)

source filters (0)

Q Filter

All field types

| name | type | format | searchable | aggregatable | excluded | controls |
|---------------|----------|--------|------------|--------------|----------|----------|
| _id | string | | ✓ | ✓ | | |
| _index | string | | ✓ | ✓ | | |
| _score | number | | | | | |
| _source | ._source | | | | | |
| _type | string | | ✓ | ✓ | | |
| env_name | string | | ✓ | | | |
| model_name | string | | ✓ | | | |
| submit_status | string | | ✓ | | | |
| tot_runtime | number | | ✓ | ✓ | | |

Kibana Console

► Alternative to CURL


Dev Tools

Console


```
1 PUT /rta-all-models3
2 {
3   "mappings": {
4     "log": {
5       "properties": {
6         "env_name": {"type": "text"},
7         "submit_status": {"type": "text"},
8         "model_name": {"type": "text"},
9         "tot_runtime": {"type": "float"}
10      }
11    }
12  }
13 }
```

```
1 {
2   "acknowledged": true,
3   "shards_acknowledged": true,
4   "index": "rta-all-models3"
5 }
```

Elasticsearch for Hadoop




Architectural overview



Architectural overview

Model Monitoring Architecture



Kafka

- ▶ Kafka is generally used for building real-time streaming
 - ▶ data pipelines that reliably get data between systems or applications
 - ▶ applications that transform or react to the streams of data
- ▶ Kafka is run as a cluster on one or more servers that can span multiple datacenters.
- ▶ The Kafka cluster stores streams of records in categories called topics.
- ▶ Each record consists of a key, a value, and a timestamp.

Kafka installation

- ▶ Download the binary: `kafka_2.12-1.0.1.tgz`
- ▶ `7z x kafka_2.12-1.0.1.tgz && 7z x kafka_2.12-1.0.1.tar`
- ▶ `sudo mv kafka_2.12-1.0.1 /opt/Kafka`

Kafka demo

- ▶ `cd /opt/Kafka/ kafka_2.12-1.0.1`
- ▶ `sudo bin/kafka-server-start.sh config/server.properties`
- ▶ `bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic testing --from-beginning`
- ▶ `bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic testing`
- ▶ `bin/kafka-topics.sh --list --zookeeper localhost:2181`
- ▶ Configure Kafka producer `connect-file-source.properties`
- ▶ Configure Kafka consumer `connect-file-sink.properties`
- ▶ `bin/connect-standalone.sh`
`config/connect-standalone.properties`
`config/connect-file-source.properties`
`config/connect-file-sink.properties`

Beats

- ▶ Seamlessly integrates with the Elastic Stack. Kafka requires a separate install.
- ▶ One way, Kafka is bidirectional
- ▶ Extensible
- ▶ Shippers: Filebeat, Metricbeat, Packetbeat, Winlogbeat, Auditbeat, Heartbeat.
- ▶ Filebeat configuration

Beats demo: Heartbeat

- ▶ elastic/stack-docker

Logstash.conf

```
input {
  heartbeat {
    interval => 5
    message  => 'Hello from Logstash 🐶'
  }
}

#input {
#  filebeat {
#    port => 5044
#  }
#}

output {
  elasticsearch {
    hosts      => [ 'elasticsearch' ]
    user       => 'elastic'
    password   => 'changeme'
  }
}
```

Logstash

- ▶ grok – better pattern matching
- ▶ <https://www.elastic.co/guide/en/logstash/current/plugins-filters-grok.html>

grok

helps parse arbitrary text and structures it.
labels instead of regex patterns.

```
USERNAME [a-zA-Z0-9_-]+
USER %{USERNAME}
INT (?[+-]?([0-9]+))
MONTH \b(?:Jan(?:uary)?|Feb(?:ruary)?|Mar(?:ch)?|Apr(?:il)?|May|Jun(?:e)?|Jul(?:y)...
DAY (?:(Mon(?:day)?|Tue(?:sday)?|Wed(?:nesday)?|Thu(?:rday)?|Fri...
```

COMBINEDAPACHELOG %{IPORHOST:clientip} %{USER:ident} %{USER:auth}
\[%{HTTPDATE:timestamp} \] "(?:%{WORD:verb} %{NOTSPACE:request}
(?: HTTP/%{NUMBER:httpversion})?|-)" %{NUMBER:response}
(?:%{NUMBER:bytes}|-) %{QS:referrer} %{QS:agent}

Data generation

- ▶ Out of the box logging
- ▶ Custom logging
- ▶ Data dictionaries
- ▶ Data types
- ▶ File types

Logging facility for Python

- ▶ This module defines functions and classes which implement a flexible event logging system for applications and libraries.
- ▶ The key benefit of having the logging API provided by a standard library module is that all Python modules can participate in logging, so your application log can include your own messages integrated with messages from third-party modules.

Logging facility for Python

The basic classes defined by the module, together with their functions, are:

- ▶ Loggers expose the interface that application code directly uses.
- ▶ Handlers send the log records (created by loggers) to the appropriate destination.
- ▶ Filters provide a finer grained facility for determining which log records to output.
- ▶ Formatters specify the layout of log records in the final output.

Logging facility for Python

Logging serves two purposes:

- ▶ Diagnostic logging records events related to the application's operation. If a user calls in to report an error, for example, the logs can be searched for context.
- ▶ Audit logging records events for business analysis. A user's transactions can be extracted and combined with other user details for reports or to optimize a business goal.

dsds_logging guide

1. git clone
`https://<username>@tools.standardbank.co.za/bitbucket/scm/datas
packages.git`
2. `sys.path.append("../python-packages/dsds")`
3. `import dsds.dsds_logging`
4. `config, logger = dsds_spark.get_config_and_logger(sys.argv)`
5. `sc, hiveContext = dsds_spark.get_contexts(config, sys.argv)`
6. `spark_main(config, logger, sc, hiveContext)`
7. `logger.info('Feature_extraction.py', 'feature_set_6',
feature_set_6.columns)`

Sample log (.txt)

```
INFO:20180119_102504:submit:is_test=False
INFO:20180119_102504:submit:username=a231384
INFO:20180119_102504:submit:sys.platform=linux2
INFO:20180119_102504:submit:os.name=posix
INFO:20180119_102504:submit:python.version=(2, 7, 13)
INFO:20180119_102504:submit:max folder age days=2
INFO:20180119_102504:submit:folders deleted=0
INFO:20180119_102504:submit:model_name=sbgm_anomaly_classification
INFO:20180119_102504:submit:conf_environment=dev
INFO:20180119_102504:submit:config.base.project=sbgm_anomaly_classification
INFO:20180119_102504:submit:config.base.team=dsds
INFO:20180119_102504:submit:config.base.environment=dev
INFO:20180119_102504:submit:config.cluster.venv=py2-spark1
INFO:20180119_102504:submit:config.cluster.is-local=false
INFO:20180119_102504:submit:config.cluster.driver-memory=16g
INFO:20180119_102504:submit:config.cluster.num-executors=60
INFO:20180119_102504:submit:config.cluster.executor-memory=13g
INFO:20180119_102504:submit:config.cluster.executor-cores=4
INFO:20180119_102504:submit:config.steps.step-1=create_uri_summary.py
INFO:20180119_102504:submit:config.steps.step-2=create_sessions.py
INFO:20180119_102504:submit:config.steps.step-3=apply_models.py
INFO:20180119_102504:submit:config.steps.step-4=fit_models.py
INFO:20180119_102504:submit:config.data.hist_location=hdfs:///dev/data/dsds/general/history_unzip/
INFO:20180119_102504:submit:config.data.nrt_location=hdfs:///dev/data/dsds/general/nrt/
INFO:20180119_102504:submit:config.data.uri-summaries=[hive][uri_summary]
INFO:20180119_102504:submit:config.data.uri-summaries-new=[hive][uri_summary_new]
INFO:20180119_102504:submit:config.data.sessions=[hive][sessions]
INFO:20180119_102504:submit:config.data.session-summary=[hive][session_summary]
INFO:20180119_102504:submit:config.data.session-summary-new=[hive][session_summary_new]
INFO:20180119_102504:submit:config.data.fitted-models=[hive][fitted_models]
INFO:20180119_102504:submit:config.data.scored-sessions=[hive][scored_sessions]
INFO:20180119_102504:submit:config.depends.local-packages=dsds
INFO:20180119_102504:submit:config.depends.local-files=[read_log_data.py]
INFO:20180119_102504:submit:config.model.earliest_date=2017-08-05
INFO:20180119_102504:submit:config.model.last_history_date=2017-09-17
INFO:20180119_102504:submit:config.model.first_nrt_date=2017-09-20
INFO:20180119_102504:submit:config.model.run_until=2017-09-21
INFO:20180119_102504:submit:config.model.session-timeout=5minutes
INFO:20180119_102504:submit:config.model.max-session-length=30minutes
INFO:20180119_102504:submit:config.model.time_between_fits=4weeks
INFO:20180119_102504:submit:config.model.fit_length=13weeks
INFO:20180119_102504:submit:config.model.cutoff_n_sessions=50
INFO:20180119_102504:submit:config.model.tree_depth=5
INFO:20180119_102504:submit:config.model.n_trees=20
INFO:20180119_102504:submit:config.model.random_sample_per_tree=100
INFO:20180119_102504:submit:config.log.logger_type=FILE
INFO:20180119_102504:submit:config.log.log_location=../../logging
INFO:20180119_102504:submit:config.log.log_level=INFO
INFO:20180119_102504:submit:config=OK
INFO:20180119_102504:submit:step-1=create_uri_summary.py
INFO:20180119_102504:submit:step-2=create_sessions.py
INFO:20180119_102504:submit:step-3=apply_models.py
```

Sample log (.json)

```
INFO:root:{'spark': {'home': None, 'version': '2.1.0.2.6.0.3-8', 'environment': {'PYTHONHASHSEED': '0'}, 'user': 'a231384', 'conf':  
[('spark.eventlog.enabled', 'true'), ('spark.yarn.historyServer.address', 'pdshnn1p.standardbank.co.za:18081'), ('spark.history.ui.port',  
'18081'), ('spark.driver.extraLibraryPath', '/usr/hdp/current/hadoop-client/lib/native:/usr/hdp/current/hadoop-client/lib/native/Linux-  
amd64-64'), ('spark.history.kerberos.keytab', '/etc/security/keytabs/spark.headless.keytab'), ('spark.executor.id', 'driver'),  
'spark.app.id', 'local-1526633937562'), ('spark.yarn.queue', 'default'), ('spark.driver.port', '40470'), ('spark.app.name', 'pyspark-  
shell'), ('spark.executor.extraLibraryPath', '/usr/hdp/current/hadoop-client/lib/native:/usr/hdp/current/hadoop-client/lib/native/Linux-  
amd64-64'), ('spark.driver.host', '10.144.164.203'), ('spark.history.kerberos.principal', 'spark-ds_hdp_prod@ZA.SBICDIRECTORY.COM'),  
'spark.history.fs.logDirectory', 'hdfs:///spark2-history/'), ('spark.sql.catalogImplementation', 'hive'), ('spark.rdd.compress', 'True'),  
'spark.history.provider', 'org.apache.spark.deploy.history.FsHistoryProvider'), ('spark.serializer.objectStreamReset', '100'),  
'spark.master', 'local[*]'), ('spark.submit.deployMode', 'client'), ('hive.metastore.warehouse.dir', 'file:/home/a231384/rta/anomaly-  
detection/sbg-dsds-fraud-anomaly-detection/helpers/digital_anomaly_detection/monitoring/spark-warehouse'), ('spark.port.maxRetries', '100'),  
'spark.eventlog.dir', 'hdfs:///spark2-history/')], 'python': {'version': '3.4'}, 'start_time': '2018-05-18T11:03:04.926190', 'ds_env':  
'\n', 'data': {'historical': 'hdfs:///dev/data/dsds/general/history_unzip', 'near_real_time': 'hdfs:///dev/data/dsds/general/nrt',  
'list_of_hdfs_files': 'list_of_hdfs_files.txt'}, 'modules': ['IPython.core.shadows', 'sklearn.linear_model', 'sys', 'pandas', 'json',  
'logging', 'builtins', 'pickle', 'subprocess', 'time', 'requests', 'pyspark', 'types', 'py4j', 're', 'atexit', 'os', 'datetime', 'builtins',  
'platform', 'random', 'numpy', 'configparser'], 'pyspark': {'submit': {'args': '\n'}}}  
WARNING:root:{}  
ERROR:root:{}  
INFO:root:{'spark': {'home': None, 'version': '2.1.0.2.6.0.3-8', 'environment': {'PYTHONHASHSEED': '0'}, 'user': 'a231384', 'conf':  
[('spark.eventlog.enabled', 'true'), ('spark.yarn.historyServer.address', 'pdshnn1p.standardbank.co.za:18081'), ('spark.history.ui.port',  
'18081'), ('spark.driver.extraLibraryPath', '/usr/hdp/current/hadoop-client/lib/native:/usr/hdp/current/hadoop-client/lib/native/Linux-  
amd64-64'), ('spark.history.kerberos.keytab', '/etc/security/keytabs/spark.headless.keytab'), ('spark.executor.id', 'driver'),  
'spark.app.id', 'local-1526633937562'), ('spark.yarn.queue', 'default'), ('spark.driver.port', '40470'), ('spark.app.name', 'pyspark-  
shell'), ('spark.executor.extraLibraryPath', '/usr/hdp/current/hadoop-client/lib/native:/usr/hdp/current/hadoop-client/lib/native/Linux-  
amd64-64'), ('spark.driver.host', '10.144.164.203'), ('spark.history.kerberos.principal', 'spark-ds_hdp_prod@ZA.SBICDIRECTORY.COM'),  
'spark.history.fs.logDirectory', 'hdfs:///spark2-history/'), ('spark.sql.catalogImplementation', 'hive'), ('spark.rdd.compress', 'True'),  
'spark.history.provider', 'org.apache.spark.deploy.history.FsHistoryProvider'), ('spark.serializer.objectStreamReset', '100'),  
'spark.master', 'local[*]'), ('spark.submit.deployMode', 'client'), ('hive.metastore.warehouse.dir', 'file:/home/a231384/rta/anomaly-  
detection/sbg-dsds-fraud-anomaly-detection/helpers/digital_anomaly_detection/monitoring/spark-warehouse'), ('spark.port.maxRetries', '100'),  
'spark.eventlog.dir', 'hdfs:///spark2-history/')], 'python': {'version': '3.4'}, 'start_time': '2018-05-18T11:03:04.926190', 'ds_env':  
'\n', 'data': {'historical': 'hdfs:///dev/data/dsds/general/history_unzip', 'near_real_time': 'hdfs:///dev/data/dsds/general/nrt',  
'list_of_hdfs_files': 'list_of_hdfs_files.txt'}, 'modules': ['IPython.core.shadows', 'sklearn.linear_model', 'sys', 'pandas', 'json',  
'logging', 'builtins', 'pickle', 'subprocess', 'time', 'requests', 'pyspark', 'types', 'py4j', 're', 'atexit', 'os', 'datetime', 'builtins',  
'platform', 'random', 'numpy', 'configparser'], 'pyspark': {'submit': {'args': '\n'}}}  
WARNING:root:{}  
ERROR:root:}
```

Monitoring

- ▶ Managing and monitoring statistical models is crucial if your organization periodically runs a large number (say, over 10) of statistical models.
- ▶ However, these issues are important even when there are just a few of them in production.

Monitoring

Common challenges include the following:

- ▶ Keeping all the input correct and fresh.
- ▶ Making sure the outputs go to the right places, in the correct formats.
- ▶ Keeping the code organized for effective updating and maintenance.
- ▶ Creating and maintaining effective documentation.
- ▶ Assessing and tracking model performance.
- ▶ Effectively (preferably automatically) deciding when to update the model.

Monitoring: all models

- ▶ Model: name.
- ▶ Environment: continuous development and integration; software and versions; hardware statistics; environment variables; run mode; current build version; source and run location; steps; extra packages and files; run command.
- ▶ Data: historical location; near real-time location; maximum folder age; logs start and end date; last history date; model last date; next run date; first NRT date.
- ▶ Results: submit status; total runtime; FLS alerts; loglines.

Monitoring: supervised models

- ▶ Statistical process control: drift detection method (DDM); early drift detection method (EDDM).
- ▶ Sequential analysis: linear four rates (true -ve, false -ve, true +ve, false +ve) – specificity, recall, precision, accuracy; Monte Carlo sampling for significance level; Bonferoni correction for correlated tests.
- ▶ Error distribution monitoring: adaptive windowing (ADWIN)

Monitoring: unsupervised models

- ▶ Clustering/novelty detection
- ▶ Feature distribution monitoring
- ▶ Model-dependent monitoring

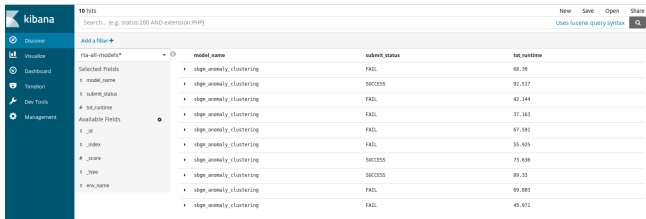
Monitoring: random forests

- ▶ Number of URI's
- ▶ Time between fits
- ▶ Samples per tree
- ▶ Model start date
- ▶ Number of sessions to score
- ▶ Previous run date
- ▶ Maximum session length
- ▶ Last URI timestamp

Monitoring: k-means

- ▶ Database name
- ▶ Results
- ▶ Alerts to FLS
- ▶ Model path
- ▶ List of features
- ▶ Clusters

Visualisation



10 hits

Search... (e.g. status:200 AND extension:PHP)

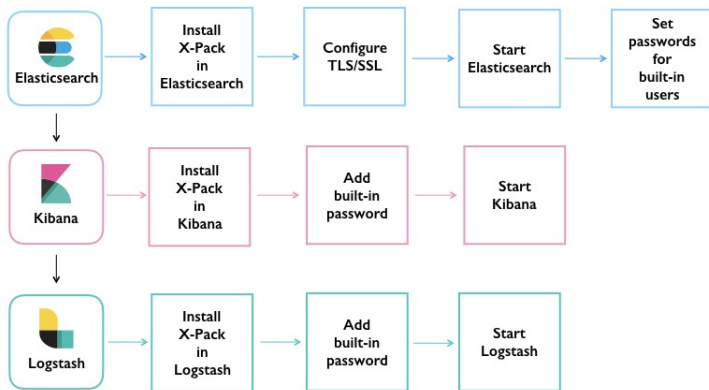
[New](#) [Save](#) [Open](#) [Share](#)

[Uses lucene query syntax](#)

[Add a filter](#)

| | model_name | submitt_status | test_runtime |
|------------------|---------------------------|----------------|--------------|
| Selected Fields | ▸ sbge_anomaly_clustering | FAIL | 88.39 |
| model_name | ▸ sbge_anomaly_clustering | SUCCESS | 92.537 |
| submitt_status | ▸ sbge_anomaly_clustering | FAIL | 42.144 |
| test_runtime | ▸ sbge_anomaly_clustering | FAIL | 37.163 |
| Available Fields | ▸ sbge_anomaly_clustering | FAIL | 67.591 |
| _id | ▸ sbge_anomaly_clustering | FAIL | 55.925 |
| _index | ▸ sbge_anomaly_clustering | SUCCESS | 73.636 |
| _score | ▸ sbge_anomaly_clustering | SUCCESS | 89.33 |
| _type | ▸ sbge_anomaly_clustering | FAIL | 69.883 |
| env_name | ▸ sbge_anomaly_clustering | FAIL | 45.971 |

X-Pack



X-Pack capabilities

- ▶ Security
- ▶ Alerting
- ▶ Monitoring
- ▶ Reporting
- ▶ Graph
- ▶ Machine learning

X-Pack: machine learning

- ▶ Time series analysis
- ▶ Anomaly detection
- ▶ <https://www.youtube.com/watch?v=n6xW6YWYgs0>

References

1. Building Intelligent Systems: A Guide to Machine Learning Engineering. [Geoff Hulten] (Apress, 2018)
2. Trends in AI, Data Science, and Big Data. [Ben Lorica] (2017)
3. Building Evolutionary Architectures. [Rebecca Parsons; Patrick Kua; Neal Ford] (O'Reilly Media, 2017)
4. 5 things you should be monitoring. [Brian Brazil] (2018)
5. The Logstash Book. [James Turnbull] (Turnbull Press, 2013)
6. Beyond the Twelve-Factor App. [Kevin Hoffman] (O'Reilly Media, 2016)
7. Logs and real-time stream processing. [Jay Kreps] (2016)
8. I Heart Logs: Apache Kafka and Real-time Data Integration. [Jay Kreps] (2015)
9. The log: The lifeblood of your data pipeline. [Kiyoto Tamura] (2015)
10. Understanding the ELK stack. [Brian Anderson; Rafał Kuć] (2016)